



memcached / memcached

[Watch](#) 593[Star](#) 5,103[Fork](#) 1,799[Code](#)[Issues 10](#)[Pull requests 5](#)[Wiki](#)[Pulse](#)[Graphs](#)Branch: [master](#) ▾[memcached / doc / protocol.txt](#)[Find file](#) [Copy path](#)

dormando document new metadump command.

c00dbc3 5 days ago

16 contributors



1086 lines (822 sloc) | 48.9 KB

[Raw](#)[Blame](#)[History](#)

```
1 Protocol
2 -----
3
4 Clients of memcached communicate with server through TCP connections.
5 (A UDP interface is also available; details are below under "UDP
6 protocol.") A given running memcached server listens on some
7 (configurable) port; clients connect to that port, send commands to
8 the server, read responses, and eventually close the connection.
9
10 There is no need to send any command to end the session. A client may
11 just close the connection at any moment it no longer needs it. Note,
12 however, that clients are encouraged to cache their connections rather
13 than reopen them every time they need to store or retrieve data. This
14 is because memcached is especially designed to work very efficiently
15 with a very large number (many hundreds, more than a thousand if
16 necessary) of open connections. Caching connections will eliminate the
17 overhead associated with establishing a TCP connection (the overhead
18 of preparing for a new connection on the server side is insignificant
19 compared to this).
20
21 There are two kinds of data sent in the memcache protocol: text lines
22 and unstructured data. Text lines are used for commands from clients
23 and responses from servers. Unstructured data is sent when a client
24 wants to store or retrieve data. The server will transmit back
25 unstructured data in exactly the same way it received it, as a byte
26 stream. The server doesn't care about byte order issues in
27 unstructured data and isn't aware of them. There are no limitations on
28 characters that may appear in unstructured data; however, the reader
29 of such data (either a client or a server) will always know, from a
30 preceding text line, the exact length of the data block being
31 transmitted.
32
33 Text lines are always terminated by \r\n. Unstructured data is _also_
34 terminated by \r\n, even though \r, \n or any other 8-bit characters
35 may also appear inside the data. Therefore, when a client retrieves
36 data from a server, it must use the length of the data block (which it
37 will be provided with) to determine where the data block ends, and not
38 the fact that \r\n follows the end of the data block, even though it
39 does.
40
41 Keys
42 -----
43
44 Data stored by memcached is identified with the help of a key. A key
45 is a text string which should uniquely identify the data for clients
46 that are interested in storing and retrieving it. Currently the
47 length limit of a key is set at 250 characters (of course, normally
48 clients wouldn't need to use such long keys); the key must not include
```

```
49 control characters or whitespace.
50
51 Commands
52 -----
53
54 There are three types of commands.
55
56 Storage commands (there are six: "set", "add", "replace", "append"
57 "prepend" and "cas") ask the server to store some data identified by a
58 key. The client sends a command line, and then a data block; after
59 that the client expects one line of response, which will indicate
60 success or failure.
61
62 Retrieval commands (there are two: "get" and "gets") ask the server to
63 retrieve data corresponding to a set of keys (one or more keys in one
64 request). The client sends a command line, which includes all the
65 requested keys; after that for each item the server finds it sends to
66 the client one response line with information about the item, and one
67 data block with the item's data; this continues until the server
68 finished with the "END" response line.
69
70 All other commands don't involve unstructured data. In all of them,
71 the client sends one command line, and expects (depending on the
72 command) either one line of response, or several lines of response
73 ending with "END" on the last line.
74
75 A command line always starts with the name of the command, followed by
76 parameters (if any) delimited by whitespace. Command names are
77 lower-case and are case-sensitive.
78
79 Expiration times
80 -----
81
82 Some commands involve a client sending some kind of expiration time
83 (relative to an item or to an operation requested by the client) to
84 the server. In all such cases, the actual value sent may either be
85 Unix time (number of seconds since January 1, 1970, as a 32-bit
86 value), or a number of seconds starting from current time. In the
87 latter case, this number of seconds may not exceed 60*60*24*30 (number
88 of seconds in 30 days); if the number sent by a client is larger than
89 that, the server will consider it to be real Unix time value rather
90 than an offset from current time.
91
92
93 Error strings
94 -----
95
96 Each command sent by a client may be answered with an error string
97 from the server. These error strings come in three types:
98
99 - "ERROR\r\n"
100
101 means the client sent a nonexistent command name.
102
103 - "CLIENT_ERROR <error>\r\n"
104
105 means some sort of client error in the input line, i.e. the input
106 doesn't conform to the protocol in some way. <error> is a
107 human-readable error string.
108
109 - "SERVER_ERROR <error>\r\n"
110
111 means some sort of server error prevents the server from carrying
112 out the command. <error> is a human-readable error string. In cases
113 of severe server errors, which make it impossible to continue
114 serving the client (this shouldn't normally happen), the server will
115 close the connection after sending the error line. This is the only
```

```
116 case in which the server closes a connection to a client.  
117  
118  
119 In the descriptions of individual commands below, these error lines  
120 are not again specifically mentioned, but clients must allow for their  
121 possibility.  
122  
123  
124 Storage commands  
125 -----  
126  
127 First, the client sends a command line which looks like this:  
128  
129 <command name> <key> <flags> <exptime> <bytes> [noreply]\r\n  
130 cas <key> <flags> <exptime> <bytes> <cas unique> [noreply]\r\n  
131  
132 - <command name> is "set", "add", "replace", "append" or "prepend"  
133  
134 "set" means "store this data".  
135  
136 "add" means "store this data, but only if the server *doesn't* already  
137 hold data for this key".  
138  
139 "replace" means "store this data, but only if the server *does*  
140 already hold data for this key".  
141  
142 "append" means "add this data to an existing key after existing data".  
143  
144 "prepend" means "add this data to an existing key before existing data".  
145  
146 The append and prepend commands do not accept flags or exptime.  
147 They update existing data portions, and ignore new flag and exptime  
148 settings.  
149  
150 "cas" is a check and set operation which means "store this data but  
151 only if no one else has updated since I last fetched it."  
152  
153 - <key> is the key under which the client asks to store the data  
154  
155 - <flags> is an arbitrary 16-bit unsigned integer (written out in  
156 decimal) that the server stores along with the data and sends back  
157 when the item is retrieved. Clients may use this as a bit field to  
158 store data-specific information; this field is opaque to the server.  
159 Note that in memcached 1.2.1 and higher, flags may be 32-bits, instead  
160 of 16, but you might want to restrict yourself to 16 bits for  
161 compatibility with older versions.  
162  
163 - <exptime> is expiration time. If it's 0, the item never expires  
164 (although it may be deleted from the cache to make place for other  
165 items). If it's non-zero (either Unix time or offset in seconds from  
166 current time), it is guaranteed that clients will not be able to  
167 retrieve this item after the expiration time arrives (measured by  
168 server time). If a negative value is given the item is immediately  
169 expired.  
170  
171 - <bytes> is the number of bytes in the data block to follow, *not*  
172 including the delimiting \r\n. <bytes> may be zero (in which case  
173 it's followed by an empty data block).  
174  
175 - <cas unique> is a unique 64-bit value of an existing entry.  
176 Clients should use the value returned from the "gets" command  
177 when issuing "cas" updates.  
178  
179 - "noreply" optional parameter instructs the server to not send the  
180 reply. NOTE: if the request line is malformed, the server can't  
181 parse "noreply" option reliably. In this case it may send the error  
182 to the client, and not reading it on the client side will break
```

```
183 things. Client should construct only valid requests.  
184  
185 After this line, the client sends the data block:  
186  
187 <data block>\r\n  
188  
189 - <data block> is a chunk of arbitrary 8-bit data of length <bytes>  
190 from the previous line.  
191  
192 After sending the command line and the data block the client awaits  
193 the reply, which may be:  
194  
195 - "STORED\r\n", to indicate success.  
196  
197 - "NOT_STORED\r\n" to indicate the data was not stored, but not  
198 because of an error. This normally means that the  
199 condition for an "add" or a "replace" command wasn't met.  
200  
201 - "EXISTS\r\n" to indicate that the item you are trying to store with  
202 a "cas" command has been modified since you last fetched it.  
203  
204 - "NOT_FOUND\r\n" to indicate that the item you are trying to store  
205 with a "cas" command did not exist.  
206  
207  
208 Retrieval command:  
209 -----  
210  
211 The retrieval commands "get" and "gets" operates like this:  
212  
213 get <key>*\r\n  
214 gets <key>*\r\n  
215  
216 - <key>* means one or more key strings separated by whitespace.  
217  
218 After this command, the client expects zero or more items, each of  
219 which is received as a text line followed by a data block. After all  
220 the items have been transmitted, the server sends the string  
221  
222 "END\r\n"  
223  
224 to indicate the end of response.  
225  
226 Each item sent by the server looks like this:  
227  
228 VALUE <key> <flags> <bytes> [<cas unique>]\r\n  
229 <data block>\r\n  
230  
231 - <key> is the key for the item being sent  
232  
233 - <flags> is the flags value set by the storage command  
234  
235 - <bytes> is the length of the data block to follow, *not* including  
236 its delimiting \r\n  
237  
238 - <cas unique> is a unique 64-bit integer that uniquely identifies  
239 this specific item.  
240  
241 - <data block> is the data for this item.  
242  
243 If some of the keys appearing in a retrieval request are not sent back  
244 by the server in the item list this means that the server does not  
245 hold items with such keys (because they were never stored, or stored  
246 but deleted to make space for more items, or expired, or explicitly  
247 deleted by a client).  
248  
249
```

```
250 Deletion
251 -----
252
253 The command "delete" allows for explicit deletion of items:
254
255 delete <key> [noreply]\r\n
256
257 - <key> is the key of the item the client wishes the server to delete
258
259 - "noreply" optional parameter instructs the server to not send the
260   reply. See the note in Storage commands regarding malformed
261   requests.
262
263 The response line to this command can be one of:
264
265 - "DELETED\r\n" to indicate success
266
267 - "NOT_FOUND\r\n" to indicate that the item with this key was not
268   found.
269
270 See the "flush_all" command below for immediate invalidation
271 of all existing items.
272
273
274 Increment/Decrement
275 -----
276
277 Commands "incr" and "decr" are used to change data for some item
278 in-place, incrementing or decrementing it. The data for the item is
279 treated as decimal representation of a 64-bit unsigned integer. If
280 the current data value does not conform to such a representation, the
281 incr/decr commands return an error (memcached <= 1.2.6 treated the
282 bogus value as if it were 0, leading to confusion). Also, the item
283 must already exist for incr/decr to work; these commands won't pretend
284 that a non-existent key exists with value 0; instead, they will fail.
285
286 The client sends the command line:
287
288 incr <key> <value> [noreply]\r\n
289
290 or
291
292 decr <key> <value> [noreply]\r\n
293
294 - <key> is the key of the item the client wishes to change
295
296 - <value> is the amount by which the client wants to increase/decrease
297   the item. It is a decimal representation of a 64-bit unsigned integer.
298
299 - "noreply" optional parameter instructs the server to not send the
300   reply. See the note in Storage commands regarding malformed
301   requests.
302
303 The response will be one of:
304
305 - "NOT_FOUND\r\n" to indicate the item with this value was not found
306
307 - <value>\r\n , where <value> is the new value of the item's data,
308   after the increment/decrement operation was carried out.
309
310 Note that underflow in the "decr" command is caught: if a client tries
311 to decrease the value below 0, the new value will be 0. Overflow in
312 the "incr" command will wrap around the 64 bit mark.
313
314 Note also that decrementing a number such that it loses length isn't
315 guaranteed to decrement its returned length. The number MAY be
316 space-padded at the end, but this is purely an implementation
```

```
317 optimization, so you also shouldn't rely on that.  
318  
319 Touch  
320 -----  
321  
322 The "touch" command is used to update the expiration time of an existing item  
323 without fetching it.  
324  
325 touch <key> <exptime> [noreply]\r\n  
326  
327 - <key> is the key of the item the client wishes the server to touch  
328  
329 - <exptime> is expiration time. Works the same as with the update commands  
330 (set/add/etc). This replaces the existing expiration time. If an existing  
331 item were to expire in 10 seconds, but then was touched with an  
332 expiration time of "20", the item would then expire in 20 seconds.  
333  
334 - "noreply" optional parameter instructs the server to not send the  
335 reply. See the note in Storage commands regarding malformed  
336 requests.  
337  
338 The response line to this command can be one of:  
339  
340 - "TOUCHED\r\n" to indicate success  
341  
342 - "NOT_FOUND\r\n" to indicate that the item with this key was not  
343 found.  
344  
345 Slabs Reassign  
346 -----  
347  
348 NOTE: This command is subject to change as of this writing.  
349  
350 The slabs reassign command is used to redistribute memory once a running  
351 instance has hit its limit. It might be desirable to have memory laid out  
352 differently than was automatically assigned after the server started.  
353  
354 slabs reassign <source class> <dest class>\r\n  
355  
356 - <source class> is an id number for the slab class to steal a page from  
357  
358 A source class id of -1 means "pick from any valid class"  
359  
360 - <dest class> is an id number for the slab class to move a page to  
361  
362 The response line could be one of:  
363  
364 - "OK" to indicate the page has been scheduled to move  
365  
366 - "BUSY [message]" to indicate a page is already being processed, try again  
367 later.  
368  
369 - "BADCLASS [message]" a bad class id was specified  
370  
371 - "NOSPARE [message]" source class has no spare pages  
372  
373 - "NOTFULL [message]" dest class must be full to move new pages to it  
374  
375 - "UNSAFE [message]" source class cannot move a page right now  
376  
377 - "SAME [message]" must specify different source/dest ids.  
378  
379 Slabs Automove  
380 -----  
381  
382 NOTE: This command is subject to change as of this writing.  
383
```

```
384 The slabs automove command enables a background thread which decides on its  
385 own when to move memory between slab classes. Its implementation and options  
386 will likely be in flux for several versions. See the wiki/mailing list for  
387 more details.  
388  
389 The automover can be enabled or disabled at runtime with this command.  
390  
391 slabs automove <0|1>  
392  
393 - 0|1|2 is the indicator on whether to enable the slabs automover or not.  
394  
395 The response should always be "OK\r\n"  
396  
397 - <0> means to set the thread on standby  
398  
399 - <1> means to return pages to a global pool when there are more than 2 pages  
400 worth of free chunks in a slab class. Pages are then re-assigned back into  
401 other classes as-needed.  
402  
403 - <2> is a highly aggressive mode which causes pages to be moved every time  
404 there is an eviction. It is not recommended to run for very long in this  
405 mode unless your access patterns are very well understood.  
406  
407 LRU_Crawler  
408 -----  
409  
410 NOTE: This command (and related commands) are subject to change as of this  
411 writing.  
412  
413 The LRU Crawler is an optional background thread which will walk from the tail  
414 toward the head of requested slab classes, actively freeing memory for expired  
415 items. This is useful if you have a mix of items with both long and short  
416 TTL's, but aren't accessed very often. This system is not required for normal  
417 usage, and can add small amounts of latency and increase CPU usage.  
418  
419 lru_crawler <enable|disable>  
420  
421 - Enable or disable the LRU Crawler background thread.  
422  
423 The response line could be one of:  
424  
425 - "OK" to indicate the crawler has been started or stopped.  
426  
427 - "ERROR [message]" something went wrong while enabling or disabling.  
428  
429 lru_crawler sleep <microseconds>  
430  
431 - The number of microseconds to sleep in between each item checked for  
432 expiration. Smaller numbers will obviously impact the system more.  
433 A value of "0" disables the sleep, "1000000" (one second) is the max.  
434  
435 The response line could be one of:  
436  
437 - "OK"  
438  
439 - "CLIENT_ERROR [message]" indicating a format or bounds issue.  
440  
441 lru_crawler tocrawl <32u>  
442  
443 - The maximum number of items to inspect in a slab class per run request. This  
444 allows you to avoid scanning all of very large slabs when it is unlikely to  
445 find items to expire.  
446  
447 The response line could be one of:  
448  
449 - "OK"
```

```
451 - "CLIENT_ERROR [message]" indicating a format or bound issue.
452
453 lru_crawler crawl <classid,classid,classid|all>
454
455 - Takes a single, or a list of, numeric classids (ie: 1,3,10). This instructs
456 the crawler to start at the tail of each of these classids and run to the
457 head. The crawler cannot be stopped or restarted until it completes the
458 previous request.
459
460 The special keyword "all" instructs it to crawl all slabs with items in
461 them.
462
463 The response line could be one of:
464
465 - "OK" to indicate successful launch.
466
467 - "BUSY [message]" to indicate the crawler is already processing a request.
468
469 - "BADCLASS [message]" to indicate an invalid class was specified.
470
471 lru_crawler metadump <classid,classid,classid|all>
472
473 - Similar in function to the above "lru_crawler crawl" command, this function
474 outputs one line for every valid item found in the matching slab classes.
475 Similar to "cachedump", but does not lock the cache and can return all
476 items, not just 1MB worth.
477
478 Lines are in "key=value key2=value2" format, with value being URI encoded
479 (ie: %20 for a space).
480
481 The exact keys available are subject to change, but will include at least:
482
483 "key", "exp" (expiration time), "la", (last access time), "cas",
484 "fetch" (if item has been fetched before).
485
486 The response line could be one of:
487
488 - "OK" to indicate successful launch.
489
490 - "BUSY [message]" to indicate the crawler is already processing a request.
491
492 - "BADCLASS [message]" to indicate an invalid class was specified.
493
494 Watchers
495 -----
496
497 Watchers are a way to connect to memcached and inspect what's going on
498 internally. This is an evolving feature so new endpoints should show up over
499 time.
500
501 watch <fetchers|mutations|evictions>
502
503 - Turn connection into a watcher. Options can be stacked and are
504 space-separated. Logs will be sent to the watcher until it disconnects.
505
506 The response line could be one of:
507
508 - "OK" to indicate the watcher is ready to send logs.
509
510 - "ERROR [message]" something went wrong while enabling.
511
512 The response format is in "key=value key2=value2" format, for easy parsing.
513 Lines are prepending with "ts=" for a timestamp and "gid=" for a global ID
514 number of the log line. Given how logs are collected internally they may be
515 printed out of order. If this is important the GID may be used to put log
516 lines back in order.
517
```

```

518 The value of keys (and potentially other things) are "URI encoded". Since most
519 keys used conform to standard ASCII, this should have no effect. For keys with
520 less standard or binary characters, "%NN"'s are inserted to represent the
521 byte, ie: "n%2Cfoo" for "n,foo".
522
523 The arguments are:
524
525 - "fetchers": Currently emits logs every time an item is fetched internally.
526 This means a "set" command would also emit an item_get log, as it checks for
527 an item before replacing it. Multigets should also emit multiple lines.
528
529 - "mutations": Currently emits logs when an item is stored in most cases.
530 Shows errors for most cases when items cannot be stored.
531
532 - "evictions": Shows some information about items as they are evicted from the
533 cache. Useful in seeing if items being evicted were actually used, and which
534 keys are getting removed.
535
536 Statistics
537 -----
538
539 The command "stats" is used to query the server about statistics it
540 maintains and other internal data. It has two forms. Without
541 arguments:
542
543 stats\r\n
544
545 it causes the server to output general-purpose statistics and
546 settings, documented below. In the other form it has some arguments:
547
548 stats <args>\r\n
549
550 Depending on <args>, various internal data is sent by the server. The
551 kinds of arguments and the data sent are not documented in this version
552 of the protocol, and are subject to change for the convenience of
553 memcache developers.
554
555
556 General-purpose statistics
557 -----
558
559 Upon receiving the "stats" command without arguments, the server sends
560 a number of lines which look like this:
561
562 STAT <name> <value>\r\n
563
564 The server terminates this list with the line
565
566 END\r\n
567
568 In each line of statistics, <name> is the name of this statistic, and
569 <value> is the data. The following is the list of all names sent in
570 response to the "stats" command, together with the type of the value
571 sent for this name, and the meaning of the value.
572
573 In the type column below, "32u" means a 32-bit unsigned integer, "64u"
574 means a 64-bit unsigned integer. '32u.32u' means two 32-bit unsigned
575 integers separated by a colon (treat this as a floating point number).
576
577 |-----+-----+-----|
578 | Name      | Type   | Meaning          |
579 |-----+-----+-----|
580 | pid       | 32u    | Process id of this server process |
581 | uptime    | 32u    | Number of secs since the server started |
582 | time      | 32u    | current UNIX time according to the server |
583 | version   | string  | Version string of this server |
584 | pointer_size | 32     | Default size of pointers on the host OS |

```

585			(generally 32 or 64)	
586	rusage_user	32u.32u	Accumulated user time for this process	
587			(seconds:microseconds)	
588	rusage_system	32u.32u	Accumulated system time for this process	
589			(seconds:microseconds)	
590	curr_items	64u	Current number of items stored	
591	total_items	64u	Total number of items stored since	
592			the server started	
593	bytes	64u	Current number of bytes used	
594			to store items	
595	curr_connections	32u	Number of open connections	
596	total_connections	32u	Total number of connections opened since	
597			the server started running	
598	rejected_connections	64u	Connns rejected in maxconns_fast mode	
599	connection_structures	32u	Number of connection structures allocated	
600			by the server	
601	reserved_fds	32u	Number of misc fds used internally	
602	cmd_get	64u	Cumulative number of retrieval reqs	
603	cmd_set	64u	Cumulative number of storage reqs	
604	cmd_flush	64u	Cumulative number of flush reqs	
605	cmd_touch	64u	Cumulative number of touch reqs	
606	get_hits	64u	Number of keys that have been requested	
607			and found present	
608	get_misses	64u	Number of items that have been requested	
609			and not found	
610	get_expired	64u	Number of items that have been requested	
611			but had already expired.	
612	get_flushed	64u	Number of items that have been requested	
613			but have been flushed via flush_all	
614	delete_misses	64u	Number of deletions reqs for missing keys	
615	delete_hits	64u	Number of deletion reqs resulting in	
616			an item being removed.	
617	incr_misses	64u	Number of incr reqs against missing keys.	
618	incr_hits	64u	Number of successful incr reqs.	
619	decr_misses	64u	Number of decr reqs against missing keys.	
620	decr_hits	64u	Number of successful decr reqs.	
621	cas_misses	64u	Number of CAS reqs against missing keys.	
622	cas_hits	64u	Number of successful CAS reqs.	
623	cas_badval	64u	Number of CAS reqs for which a key was	
624			found, but the CAS value did not match.	
625	touch_hits	64u	Numer of keys that have been touched with	
626			a new expiration time	
627	touch_misses	64u	Numer of items that have been touched and	
628			not found	
629	auth_cmds	64u	Number of authentication commands	
630			handled, success or failure.	
631	auth_errors	64u	Number of failed authentications.	
632	idle_kicks	64u	Number of connections closed due to	
633			reaching their idle timeout.	
634	evictions	64u	Number of valid items removed from cache	
635			to free memory for new items	
636	reclaimed	64u	Number of times an entry was stored using	
637			memory from an expired entry	
638	bytes_read	64u	Total number of bytes read by this server	
639			from network	
640	bytes_written	64u	Total number of bytes sent by this server	
641			to network	
642	limit_maxbytes	size_t	Number of bytes this server is allowed to	
643			use for storage.	
644	accepting_conns	bool	Whether or not server is accepting conns	
645	listen_disabled_num	64u	Number of times server has stopped	
646			accepting new connections (maxconns).	
647	time_in_listen_disabled_us			
648		64u	Number of microseconds in maxconns.	
649	threads	32u	Number of worker threads requested.	
650			(see doc/threads.txt)	
651	conn_yields	64u	Number of times any connection yielded to	

652		another due to hitting the -R limit.	
653	hash_power_level	32u	Current size multiplier for hash table
654	hash_bytes	64u	Bytes currently used by hash tables
655	hash_is_expanding	bool	Indicates if the hash table is being
656		grown to a new size	
657	expired_unfetched	64u	Items pulled from LRU that were never
658		touched by get/incr/append/etc before	
659		expiring	
660	evicted_unfetched	64u	Items evicted from LRU that were never
661		touched by get/incr/append/etc.	
662	slab_reassign_running	bool	If a slab page is being moved
663	slabs_moved	64u	Total slab pages moved
664	crawler_reclaimed	64u	Total items freed by LRU Crawler
665	crawler_items-checked	64u	Total items examined by LRU Crawler
666	lru_tail_reflocked	64u	Times LRU tail was found with active ref.
667		Items can be evicted to avoid OOM errors.	
668	moves_to_cold	64u	Items moved from HOT/WARM to COLD LRU's
669	moves_to_warm	64u	Items moved from COLD to WARM LRU
670	moves_within_lru	64u	Items reshuffled within HOT or WARM LRU's
671	direct_reclaims	64u	Times worker threads had to directly
672		reclaim or evict items.	
673	lru_crawler_starts	64u	Times an LRU crawler was started
674	lru_maintainer_juggles		
675		64u	Number of times the LRU bg thread woke up
676	slab_global_page_pool	32u	Slab pages returned to global pool for
677		reassignment to other slab classes.	
678	slab_reassign_rescues	64u	Items rescued from eviction in page move
679	slab_reassign_evictions_nomem		
680		64u	Valid items evicted during a page move
681		(due to no free memory in slab)	
682	slab_reassign_chunk_rescues		
683		64u	Individual sections of an item rescued
684		during a page move.	
685	slab_reassign_inline_reclaim		
686		64u	Internal stat counter for when the page
687		mover clears memory from the chunk	
688		freelist when it wasn't expecting to.	
689	slab_reassign_busy_items		
690		64u	Items busy during page move, requiring a
691		retry before page can be moved.	
692	log_worker_dropped	64u	Logs a worker never wrote due to full buf
693	log_worker_written	64u	Logs written by a worker, to be picked up
694	log_watcher_skipped	64u	Logs not sent to slow watchers.
695	log_watcher_sent	64u	Logs written to watchers.
696	-----+-----+-----		

697

698 Settings statistics

699 -----

700 CAVEAT: This section describes statistics which are subject to change in the
701 future.

702

703 The "stats" command with the argument of "settings" returns details of
704 the settings of the running memcached. This is primarily made up of
705 the results of processing commandline options.

706

707 Note that these are not guaranteed to return in any specific order and
708 this list may not be exhaustive. Otherwise, this returns like any
709 other stats command.

710

711	-----+-----+-----		
712	Name	Type	Meaning
713	-----+-----+-----		
714	maxbytes	size_t	Maximum number of bytes allows in this cache
715	maxconn	32	Maximum number of clients allowed.
716	tcpport	32	TCP listen port.
717	udpport	32	UDP listen port.
718	inter	string	Listen interface.

```

719 | verbosity      | 32      | 0 = none, 1 = some, 2 = lots          |
720 | oldest         | 32u     | Age of the oldest honored object.   |
721 | evictions      | on/off   | When off, LRU evictions are disabled. |
722 | domain_socket  | string   | Path to the domain socket (if any). |
723 | umask          | 32 (oct) | umask for the creation of the domain socket. |
724 | growth_factor  | float    | Chunk size growth factor.          |
725 | chunk_size     | 32       | Minimum space allocated for key+value+flags. |
726 | num_threads    | 32       | Number of threads (including dispatch). |
727 | stat_key_prefix| char     | Stats prefix separator character.   |
728 | detail_enabled | bool     | If yes, stats detail is enabled.    |
729 | reqs_per_event | 32       | Max num IO ops processed within an event. |
730 | cas_enabled    | bool     | When no, CAS is not enabled for this server. |
731 | tcp_backlog    | 32       | TCP listen backlog.                |
732 | auth_enabled_sasl| yes/no   | SASL auth requested and enabled.   |
733 | item_size_max  | size_t   | maximum item size.               |
734 | maxconns_fast | bool     | If fast disconnects are enabled.   |
735 | hashpower_init | 32       | Starting size multiplier for hash table. |
736 | slab_reassign  | bool     | Whether slab page reassignment is allowed. |
737 | slab_automove  | bool     | Whether slab page automover is enabled. |
738 | slab_chunk_max | 32       | Max slab class size (avoid unless necessary) |
739 | hash_algorithm | char     | Hash table algorithm in use.      |
740 | lru_crawler    | bool     | Whether the LRU crawler is enabled. |
741 | lru_crawler_sleep | 32     | Microseconds to sleep between LRU crawls |
742 | lru_crawler_tocrawl |
743 |                   | 32u     | Max items to crawl per slab per run |
744 | lru_maintainer_thread |
745 |                   | bool     | Split LRU mode and background threads |
746 | hot_lru_pct    | 32       | Pct of slab memory reserved for HOT LRU |
747 | warm_lru_pct   | 32       | Pct of slab memory reserved for WARM LRU |
748 | expirezero_does_not_evict |
749 |                   | bool     | If yes, items with 0 exptime cannot evict |
750 | idle_time      | 0        | Drop connections that are idle this many |
751 |                   |          | seconds (0 disables)                 |
752 | watcher_logbuf_size |
753 |                   | 32u     | Size of internal (not socket) write buffer |
754 |                   |          | per active watcher connected.           |
755 | worker_logbuf_size| 32u   | Size of internal per-worker-thread buffer |
756 |                   |          | which the background thread reads from. |
757 | track_sizes     | bool     | If yes, a "stats sizes" histogram is being |
758 |                   |          | dynamically tracked.                  |
759 |-----+-----+-----|
760
761
762 Item statistics
763 -----
764 CAVEAT: This section describes statistics which are subject to change in the
765 future.
766
767 The "stats" command with the argument of "items" returns information about
768 item storage per slab class. The data is returned in the format:
769
770 STAT items:<slabclass>:<stat> <value>\r\n
771
772 The server terminates this list with the line
773
774 END\r\n
775
776 The slabclass aligns with class ids used by the "stats slabs" command. Where
777 "stats slabs" describes size and memory usage, "stats items" shows higher
778 level information.
779
780 The following item values are defined as of writing.
781
782 Name          Meaning
783 -----
784 number        Number of items presently stored in this class. Expired
785 items are not automatically excluded.

```

```
786 number_hot Number of items presently stored in the HOT LRU.  
787 number_warm Number of items presently stored in the WARM LRU.  
788 number_cold Number of items presently stored in the COLD LRU.  
789 number_noexp Number of items presently stored in the NOEXP class.  
790 age Age of the oldest item in the LRU.  
791 evicted Number of times an item had to be evicted from the LRU  
792 before it expired.  
793 evicted_nonzero Number of times an item which had an explicit expire  
794 time set had to be evicted from the LRU before it  
795 expired.  
796 evicted_time Seconds since the last access for the most recent item  
797 evicted from this class. Use this to judge how  
798 recently active your evicted data is.  
799 outofmemory Number of times the underlying slab class was unable to  
800 store a new item. This means you are running with -M or  
801 an eviction failed.  
802 tailrepairs Number of times we self-healed a slab with a refcount  
803 leak. If this counter is increasing a lot, please  
804 report your situation to the developers.  
805 reclaimed Number of times an entry was stored using memory from  
806 an expired entry.  
807 expired_unfetched Number of expired items reclaimed from the LRU which  
808 were never touched after being set.  
809 evicted_unfetched Number of valid items evicted from the LRU which were  
810 never touched after being set.  
811 crawler_reclaimed Number of items freed by the LRU Crawler.  
812 lrutail_reflocked Number of items found to be refcount locked in the  
813 LRU tail.  
814 moves_to_cold Number of items moved from HOT or WARM into COLD.  
815 moves_to_warm Number of items moved from COLD to WARM.  
816 moves_within_lru Number of times active items were bumped within  
817 HOT or WARM.  
818 direct_reclaims Number of times worker threads had to directly pull LRU  
819 tails to find memory for a new item.  
820  
821 Note this will only display information about slabs which exist, so an empty  
822 cache will return an empty set.  
823  
824 Item size statistics  
-----  
825 CAVEAT: This section describes statistics which are subject to change in the  
826 future.  
827 The "stats" command with the argument of "sizes" returns information about the  
828 general size and count of all items stored in the cache.  
829 WARNING: In versions prior to 1.4.27 this command causes the cache server to  
830 lock while it iterates the items. 1.4.27 and greater are safe.  
831  
832 The data is returned in the following format:  
833  
834 STAT <size> <count>\r\n  
835 The server terminates this list with the line  
836  
837 END\r\n  
838  
839 'size' is an approximate size of the item, within 32 bytes.  
840 'count' is the amount of items that exist within that 32-byte range.  
841  
842 This is essentially a display of all of your items if there was a slab class  
843 for every 32 bytes. You can use this to determine if adjusting the slab growth  
844 factor would save memory overhead. For example: generating more classes in the  
845 lower range could allow items to fit more snugly into their slab classes, if  
846 most of your items are less than 200 bytes in size.  
847  
848 In 1.4.27 and after, this feature must be manually enabled.
```

```

853
854 A "stats" command with the argument of "sizes_enable" will enable the
855 histogram at runtime. This has a small overhead to every store or delete
856 operation. If you don't want to incur this, leave it off.
857
858 A "stats" command with the argument of "sizes_disable" will disable the
859 histogram.
860
861 It can also be enabled at starttime with "-o track_sizes"
862
863 If disabled, "stats sizes" will return:
864
865 STAT sizes_status disabled\r\n
866
867 "stats sizes_enable" will return:
868
869 STAT sizes_status enabled\r\n
870
871 "stats sizes_disable" will return:
872
873 STAT sizes_status disabled\r\n
874
875 If an error happens, it will return:
876
877 STAT sizes_status error\r\n
878 STAT sizes_error [error_message]\r\n
879
880 CAVEAT: If CAS support is disabled, you cannot enable/disable this feature at
881 runtime.
882
883 Slab statistics
884 -----
885 CAVEAT: This section describes statistics which are subject to change in the
886 future.
887
888 The "stats" command with the argument of "slabs" returns information about
889 each of the slabs created by memcached during runtime. This includes per-slab
890 information along with some totals. The data is returned in the format:
891
892 STAT <slabclass>:<stat> <value>\r\n
893 STAT <stat> <value>\r\n
894
895 The server terminates this list with the line
896
897 END\r\n
898
899 |-----+-----|
900 | Name      | Meaning
901 |-----+-----|
902 | chunk_size | The amount of space each chunk uses. One item will use
903 |           | one chunk of the appropriate size.
904 | chunks_per_page | How many chunks exist within one page. A page by
905 |           | default is less than or equal to one megabyte in size.
906 |           | Slabs are allocated by page, then broken into chunks.
907 | total_pages | Total number of pages allocated to the slab class.
908 | total_chunks | Total number of chunks allocated to the slab class.
909 | get_hits | Total number of get requests serviced by this class.
910 | cmd_set | Total number of set requests storing data in this class.
911 | delete_hits | Total number of successful deletes from this class.
912 | incr_hits | Total number of incr operations modifying this class.
913 | decr_hits | Total number of decr operations modifying this class.
914 | cas_hits | Total number of CAS commands modifying this class.
915 | cas_badval | Total number of CAS commands that failed to modify a
916 |           | value due to a bad CAS id.
917 | touch_hits | Total number of touches serviced by this class.
918 | used_chunks | How many chunks have been allocated to items.
919 | free_chunks | Chunks not yet allocated to items, or freed via delete.

```

```

920 | free_chunks_end | Number of free chunks at the end of the last allocated | |
921 | | page. | |
922 | mem_requested | Number of bytes requested to be stored in this slab[*]. | |
923 | active_slabs | Total number of slab classes allocated. | |
924 | total_malloced | Total amount of memory allocated to slab pages. | |
925 |-----+-----| |
926
927 * Items are stored in a slab that is the same size or larger than the
928 item. mem_requested shows the size of all items within a
929 slab. (total_chunks * chunk_size) - mem_requested shows memory
930 wasted in a slab class. If you see a lot of waste, consider tuning
931 the slab factor.
932
933
934 Connection statistics
935 -----
936 The "stats" command with the argument of "conns" returns information
937 about currently active connections and about sockets that are listening
938 for new connections. The data is returned in the format:
939
940 STAT <file descriptor>:<stat> <value>\r\n
941
942 The server terminates this list with the line
943
944 END\r\n
945
946 The following "stat" keywords may be present:
947
948 |-----+-----|
949 | Name | Meaning | |
950 |-----+-----|
951 | addr | The address of the remote side. For listening | |
952 | | | sockets this is the listen address. Note that some | |
953 | | | socket types (such as UNIX-domain) don't have | |
954 | | | meaningful remote addresses. | |
955 | state | The current state of the connection. See below. | |
956 | secs_since_last_cmd | The number of seconds since the most recently | |
957 | | | issued command on the connection. This measures | |
958 | | | the time since the start of the command, so if | |
959 | | | "state" indicates a command is currently executing, | |
960 | | | this will be the number of seconds the current | |
961 | | | command has been running. | |
962 |-----+-----|
963
964 The value of the "state" stat may be one of the following:
965
966 |-----+-----|
967 | Name | Meaning | |
968 |-----+-----|
969 | conn_closing | Shutting down the connection. | |
970 | conn_listening | Listening for new connections or a new UDP request. | |
971 | conn_mwrite | Writing a complex response, e.g., to a "get" command. | |
972 | conn_new_cmd | Connection is being prepared to accept a new command. | |
973 | conn_nread | Reading extended data, typically for a command such as | |
974 | | | "set" or "put". | |
975 | conn_parse_cmd | The server has received a command and is in the middle | |
976 | | | of parsing it or executing it. | |
977 | conn_read | Reading newly-arrived command data. | |
978 | conn_swallow | Discarding excess input, e.g., after an error has | |
979 | | | occurred. | |
980 | conn_waiting | A partial command has been received and the server is | |
981 | | | waiting for the rest of it to arrive (note the difference | |
982 | | | between this and conn_nread). | |
983 | conn_write | Writing a simple response (anything that doesn't involve | |
984 | | | sending back multiple lines of response data). | |
985 |-----+-----|
986

```

```
987
988
989 Other commands
990 -----
991
992 "flush_all" is a command with an optional numeric argument. It always
993 succeeds, and the server sends "OK\r\n" in response (unless "noreply"
994 is given as the last parameter). Its effect is to invalidate all
995 existing items immediately (by default) or after the expiration
996 specified. After invalidation none of the items will be returned in
997 response to a retrieval command (unless it's stored again under the
998 same key *after* flush_all has invalidated the items). flush_all
999 doesn't actually free all the memory taken up by existing items; that
1000 will happen gradually as new items are stored. The most precise
1001 definition of what flush_all does is the following: it causes all
1002 items whose update time is earlier than the time at which flush_all
1003 was set to be executed to be ignored for retrieval purposes.
1004
1005 The intent of flush_all with a delay, was that in a setting where you
1006 have a pool of memcached servers, and you need to flush all content,
1007 you have the option of not resetting all memcached servers at the
1008 same time (which could e.g. cause a spike in database load with all
1009 clients suddenly needing to recreate content that would otherwise
1010 have been found in the memcached daemon.
1011
1012 The delay option allows you to have them reset in e.g. 10 second
1013 intervals (by passing 0 to the first, 10 to the second, 20 to the
1014 third, etc. etc.).
1015
1016 "cache_memlimit" is a command with a numeric argument. This allows runtime
1017 adjustments of the cache memory limit. It returns "OK\r\n" or an error (unless
1018 "noreply" is given as the last parameter). If the new memory limit is higher
1019 than the old one, the server may start requesting more memory from the OS. If
1020 the limit is lower, and slabs_reassign+automove are enabled, free memory may
1021 be released back to the OS asynchronously.
1022
1023 "version" is a command with no arguments:
1024
1025 version\r\n
1026
1027 In response, the server sends
1028
1029 "VERSION <version>\r\n", where <version> is the version string for the
1030 server.
1031
1032 "verbosity" is a command with a numeric argument. It always succeeds,
1033 and the server sends "OK\r\n" in response (unless "noreply" is given
1034 as the last parameter). Its effect is to set the verbosity level of
1035 the logging output.
1036
1037 "quit" is a command with no arguments:
1038
1039 quit\r\n
1040
1041 Upon receiving this command, the server closes the
1042 connection. However, the client may also simply close the connection
1043 when it no longer needs it, without issuing this command.
1044
1045
1046 UDP protocol
1047 -----
1048
1049 For very large installations where the number of clients is high enough
1050 that the number of TCP connections causes scaling difficulties, there is
1051 also a UDP-based interface. The UDP interface does not provide guaranteed
1052 delivery, so should only be used for operations that aren't required to
1053 succeed; typically it is used for "get" requests where a missing or
```

```
1054 incomplete response can simply be treated as a cache miss.  
1055  
1056 Each UDP datagram contains a simple frame header, followed by data in the  
1057 same format as the TCP protocol described above. In the current  
1058 implementation, requests must be contained in a single UDP datagram, but  
1059 responses may span several datagrams. (The only common requests that would  
1060 span multiple datagrams are huge multi-key "get" requests and "set"  
1061 requests, both of which are more suitable to TCP transport for reliability  
1062 reasons anyway.)  
1063  
1064 The frame header is 8 bytes long, as follows (all values are 16-bit integers  
1065 in network byte order, high byte first):  
1066  
1067 0-1 Request ID  
1068 2-3 Sequence number  
1069 4-5 Total number of datagrams in this message  
1070 6-7 Reserved for future use; must be 0  
1071  
1072 The request ID is supplied by the client. Typically it will be a  
1073 monotonically increasing value starting from a random seed, but the client  
1074 is free to use whatever request IDs it likes. The server's response will  
1075 contain the same ID as the incoming request. The client uses the request ID  
1076 to differentiate between responses to outstanding requests if there are  
1077 several pending from the same server; any datagrams with an unknown request  
1078 ID are probably delayed responses to an earlier request and should be  
1079 discarded.  
1080  
1081 The sequence number ranges from 0 to n-1, where n is the total number of  
1082 datagrams in the message. The client should concatenate the payloads of the  
1083 datagrams for a given response in sequence number order; the resulting byte  
1084 stream will contain a complete response in the same format as the TCP  
1085 protocol (including terminating \r\n sequences).
```

